

# Differential Privacy in Reinforcement Learning

Waël Doulazmi and Ambroise Odonnat

Master MVA, ENS Paris-Saclay  
*name.surname@ens-paris-saclay.fr*

Reinforcement Learning  
February 6, 2023

# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning
- 4 Local Differential Privacy
- 5 Discussion

# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning
- 4 Local Differential Privacy
- 5 Discussion

# Reinforcement Learning

## Definition

Reinforcement Learning is about **agents** learning in an interactive environment using **feedbacks** from their previous actions.

## Principle [Sutton and Barto, 2018]

- State space  $\mathcal{S}$ , action space  $\mathcal{A}$ , policy  $\pi : \mathcal{S} \mapsto \mathcal{A}$
- $P(s'|s, a)$  transition probability, reward  $r : \mathcal{S} \times \mathcal{A} \mapsto \mathbb{R}$
- Markov Decision Process (MDP)  $M = (\mathcal{S}, \mathcal{A}, R, P)$
- $V^\pi$  value function,  $Q^\pi$  action-value function
- Observe  $s_t \in \mathcal{S}$ , select  $a_t \in \mathcal{A}$ , receive reward  $r_t = R(s_t, a_t)$

The goal is to **maximize** the **expected cumulative rewards**

# Reinforcement Learning

## Definition

Reinforcement Learning is about **agents** learning in an interactive environment using **feedbacks** from their previous actions.




Figure 1: RL loop ([source](#))

# $(\epsilon, \delta)$ –Differential Privacy

## Principle [Dwork et al., 2006]

Mechanism to guarantee that it is **statistically hard** to **infer information** about the environment by observing learned policies.

## Notations

- Privacy budget  $\epsilon$ , probability of error  $\delta$
- The **higher** the value of  $\epsilon$ , the **lower** the privacy
- $\delta$  takes into account bad events

# Central Differential Privacy (CDP)

## Definition [Dwork et al., 2006]

Mechanism  $\mathcal{M} : D \mapsto R$  satisfies  $(\epsilon, \delta)$ -CDP if for any two adjacent inputs  $d, d' \in D$  and for any subset of outputs  $S \subseteq R$ , it holds that:

$$\mathbb{P}[\mathcal{M}(d) \in S] \leq \exp(\epsilon) \mathbb{P}[\mathcal{M}(d') \in S] + \delta \quad (1)$$




Figure 2: Central Differential Privacy ([source](#))

# Local Differential Privacy (LDP)

## Definition [Duchi et al., 2013]

Mechanism  $\mathcal{M}$  satisfies  $(\epsilon, \delta)$ -LDP if and only if for all users  $u, u' \in \mathcal{U}$ , trajectories  $(X_u, X_{u'}) \in \mathcal{X}_u \times \mathcal{X}_{u'}$  and all  $\mathcal{O} \in \{\mathcal{M}(\mathcal{X}_u) | u \in \mathcal{U}\}$ :

$$\mathbb{P}[\mathcal{M}(X_u) \in \mathcal{O}] \leq \exp(\epsilon) \mathbb{P}[\mathcal{M}(X_{u'}) \in \mathcal{O}] + \delta \quad (2)$$




Figure 3: Local Differential Privacy ([source](#))



## Related Work

Several recent papers study private RL algorithms for:

- Tabular RL [Garcelon et al., 2021]
- Bandits [Gajane et al., 2018, Chen et al., 2018]

→ Lack of private deep RL algorithms with continuous state spaces

## Contributions

- CDP versions of REINFORCE and DQN using DP-SGD
- LDP versions of REINFORCE and DQN
- Study the impact of privacy on the learning process

# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning
- 4 Local Differential Privacy
- 5 Discussion

# Cartpole

- Control problem:  $\mathcal{S}$  continuous,  $\mathcal{A}$  discrete
- Classic benchmark for RL algorithms
- 4-dimensional state  $\rightarrow$  low computational cost

Figure 4: Cartpole (OpenAI Gym)

## MDP

- $\mathcal{A} = \{0, 1\}$  (push to the left / right)
- $\mathcal{S} \subset \mathbb{R}^4$  (position / speed of the cart and the pole)
- $r(s, a) = 1$  if the pole didn't fail, 0 else

# Cartpole

## MDP

- $\mathcal{A} = \{0, 1\}$  (push to the left / right)
- $\mathcal{S} \subset \mathbb{R}^4$  (position / speed of the cart and the pole)
- $r(s, a) = 1$  if the pole didn't fail, 0 else

## Episode

- $\tau = (s_0, a_0, r_0, s_1, \dots, s_{H-1}, a_{H-1}, r_{H-1}, s_H)$
- $s_0 \sim \mu_0$  (uniform)
- Termination when:
  - ★ Pole fails
  - ★ Reach 500 steps

# Cartpole

## MDP

- $\mathcal{A} = \{0, 1\}$  (push to the left / right)
- $\mathcal{S} \subset \mathbb{R}^4$  (position / speed of the cart and the pole)
- $r(s, a) = 1$  if the pole didn't fail, 0 else

## Episode

- $\tau = (s_0, a_0, r_0, s_1, \dots, s_{H-1}, a_{H-1}, r_{H-1}, s_H)$
- $s_0 \sim \mu_0$  (uniform)
- Termination when:
  - ★ Pole fails
  - ★ Reach 500 steps

→ Solved when average reward for 25 consecutive trials  $\geq 475$

## Policy Approximation

$$\pi \in \mathcal{F}_\pi = \left\{ \pi_\theta : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1], \forall \mathbf{s} \in \mathcal{S}, \sum_{\mathbf{a} \in \mathcal{A}} \pi_\theta(\mathbf{s}, \mathbf{a}) = 1, \theta \in \mathbb{R}^d \right\}$$

## Policy Approximation

$$\pi \in \mathcal{F}_\pi = \left\{ \pi_\theta : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1], \forall \mathbf{s} \in \mathcal{S}, \sum_{\mathbf{a} \in \mathcal{A}} \pi_\theta(\mathbf{s}, \mathbf{a}) = 1, \theta \in \mathbb{R}^d \right\}$$

## Criterion - Policy performance

$$\max_{\theta \in \mathbb{R}^d} J(\pi_\theta), \quad J(\pi_\theta) = \mathbb{E}_\tau \left[ \sum_{t=0}^{H-1} \gamma^t r_t \mid \mathbf{a}_t \sim \pi_\theta(\mathbf{s}_t), \mathbf{s}_0 \sim \mu_0 \right]$$



## Policy Approximation

$$\pi \in \mathcal{F}_\pi = \left\{ \pi_\theta : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1], \forall \mathbf{s} \in \mathcal{S}, \sum_{\mathbf{a} \in \mathcal{A}} \pi_\theta(\mathbf{s}, \mathbf{a}) = 1, \theta \in \mathbb{R}^d \right\}$$

## Criterion - Policy performance

$$\max_{\theta \in \mathbb{R}^d} J(\pi_\theta), \quad J(\pi_\theta) = \mathbb{E}_\tau \left[ \sum_{t=0}^{H-1} \gamma^t r_t \mid \mathbf{a}_t \sim \pi_\theta(\mathbf{s}_t), \mathbf{s}_0 \sim \mu_0 \right]$$

## Policy Gradient

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_\tau \left[ \left( \sum_{t=0}^{H-1} \nabla_\theta \log(\pi_\theta(\mathbf{s}_t, \mathbf{a}_t)) \right) \left( \sum_{t=0}^{H-1} \gamma^t r_t \right) \right]$$

## Implementation

- 2 layers perceptron,  $h = 128$ , dropout = 0.5
- Estimate gradient with one episode




Figure 5: Convergence of REINFORCE on Cartpole

## Value Approximation

$$Q \in \mathcal{F}_Q = \{Q_\theta : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}, \theta \in \mathbb{R}^d\}$$

## Value Approximation

$$Q \in \mathcal{F}_Q = \{Q_\theta : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}, \theta \in \mathbb{R}^d\}$$

## Criterion - Temporal Difference

$$\min_{\theta \in \mathbb{R}^d} \mathcal{L}(\theta), \quad \mathcal{L}(\theta) = \|Q_\theta(s, a) - (r + \gamma \max_{a'} Q_\theta(s', a'))\|_2$$

## Value Approximation

$$Q \in \mathcal{F}_Q = \{Q_\theta : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}, \theta \in \mathbb{R}^d\}$$

## Criterion - Temporal Difference

$$\min_{\theta \in \mathbb{R}^d} \mathcal{L}(\theta), \quad \mathcal{L}(\theta) = \|Q_\theta(s, a) - (r + \gamma \max_{a'} Q_\theta(s', a'))\|_2$$

## Deep Q-Network

- Generate a transition using  $\pi_{Q_\theta}^\epsilon$  ( $\epsilon$ -greedy), add to buffer  $\mathcal{D}$
- Sample batch  $\mathcal{B} \subset \mathcal{D}$ , perform SGD

## Implementation

- 3 layers perceptron,  $h = 128$
- $|\mathcal{B}| = 128$  transitions




Figure 6: Convergence of DQN on Cartpole

# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning**
- 4 Local Differential Privacy
- 5 Discussion

## Principle [Abadi et al., 2016]

- Differentially private implementation of SGD
- AWGN noise on gradients of each sample during training
- Notable efforts to keep fast and parallel computation

## Challenges

- Adapt noise strategy for REINFORCE and DQN
- Custom backpropagation for DQN



**Input:** Examples  $\{x_1, \dots, x_N\}$ , loss function  $\mathcal{L}(\theta) = \frac{1}{N} \sum_i \mathcal{L}(\theta, x_i)$ . Parameters: learning rate  $\eta_t$ , noise scale  $\sigma$ , group size  $L$ , gradient norm bound  $C$ .

**Initialize**  $\theta_0$  randomly

**for**  $t \in [T]$  **do**

    Take a random sample  $L_t$  with sampling probability  $L/N$

**Compute gradient**

        For each  $i \in L_t$ , compute  $\mathbf{g}_t(x_i) \leftarrow \nabla_{\theta_t} \mathcal{L}(\theta_t, x_i)$

**Clip gradient**

$\bar{\mathbf{g}}_t(x_i) \leftarrow \mathbf{g}_t(x_i) / \max(1, \frac{\|\mathbf{g}_t(x_i)\|_2}{C})$

**Add noise**

$\tilde{\mathbf{g}}_t \leftarrow \frac{1}{L} (\sum_i \bar{\mathbf{g}}_t(x_i) + \mathcal{N}(0, \sigma^2 C^2 \mathbf{I}))$

**Descent**

$\theta_{t+1} \leftarrow \theta_t - \eta_t \tilde{\mathbf{g}}_t$

**Output**  $\theta_T$  and compute the overall privacy cost  $(\epsilon, \delta)$  using a privacy accounting method.

Figure 7: DP-SGD

# REINFORCE - Add noise on gradients

---

```
optimizer.zero_grad()
returns = returns.detach()
loss = - (returns * log_prob_actions).sum()
loss.backward()  # Fills p.grad
# Add noise to the gradients
for p in policy.parameters():
    clip_grad_norm_(p, max_grad_norm)

    # Gaussian noise mechanism
    p.grad += torch.normal(mean=0.,
                           std=sigma * max_grad_norm,
                           size=p.size())

# Optimize model
optimizer.step()
```

# DQN - Custom Backpropagation

---

```
# Functional to compute gradient for each sample of a batch
ft_compute_grad = grad(compute_loss_stateless_model)
ft_compute_sample_grad = vmap(ft_compute_grad,
                               in_dims=(None, None, 0, 0,
                                         0, None, None, None))

# Obtain functional version of model, params and buffers
fmodel, params, buffers = make_functional_with_buffers(policy_net)

# Recover gradients of all params for each sample of a batch
sample_grads = ft_compute_sample_grad(params, buffers, state_batch,
                                       action_batch, expected_Q_values,
                                       max_grad_norm, fmodel, criterion)
```

---

# DQN - Add noise on gradients

```
for i, param in enumerate(policy_net.parameters()):


    # Clip inplace gradients seperately
    # for each sample of the batch
    sample_grad = sample_grads[i]
    torch.nn.utils.clip_grad_norm_(sample_grad,
                                    max_norm=max_grad_norm)

    # Aggregate gradients to have a unique grad for the batch
    param.grad = torch.mean(sample_grad, dim=0)

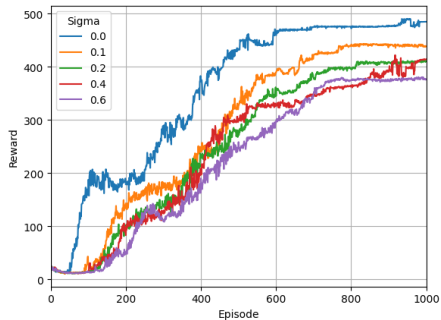
    # Add noise
    noise = torch.normal(mean=0, std=sigma*max_grad_norm,
                          size=param.size()).to(device)
    param.grad += noise/batch_size # Gaussian noise

# Optimize model
optimizer.step()
```

# CDP - Results




(a) REINFORCE




(b) DQN

Figure 8: Impact of noise level on the average agent's reward

# CDP - Results



(a) REINFORCE



(b) DQN

Figure 9: Impact of noise level on the convergence time

# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning
- 4 Local Differential Privacy**
- 5 Discussion

# LDP for Reinforcement Learning

## Method

- Disturb observations before they are seen by the agent
- Observations are the states and rewards
- Add gaussian noise on the states observed by agents




Figure 10: Local Differential Privacy ([source](#))



# LDP - Noise calibration

## Cartpole

- Reward signal is discrete and indicates episode end
- $s \in \mathbb{R}^4$ , all its component have a different range

→ Noise  $\mathcal{N}(0_4, \sigma \times IQR)$






Figure 11: Evolution of  $s$  during one training

# LDP - Results




(a) REINFORCE




(b) DQN

Figure 12: Impact of noise level on the average agent's reward

# LDP - Results



(a) REINFORCE




(b) DQN

Figure 13: Impact of noise level on the convergence time


# Overview

- 1 Background
- 2 Experimental Setup
- 3 DP-SGD for Reinforcement Learning
- 4 Local Differential Privacy
- 5 Discussion**


# Summary




(a) DP-SGD REINFORCE



(b) DP-SGD DQN



(c) REINFORCE with LDP



(d) DQN with LDP

Figure 14: Impact of noise level on the median agent's reward

# Conclusion

## Interesting results

- Implemented DP versions of two deep RL algorithms
- Proof of concept on the Cartpole environment
- Privacy vs Sample-efficiency trade-off
- Algorithms are able to converge with reasonable noise levels

## Further work

- Need for theoretical guarantees on privacy
- Explore scalability on larger environments

Thanks for your attention !

# References



Sutton and Barto (2018)

Reinforcement Learning: An Introduction

*The MIT Press*



Dwork et al. (2006)

Calibrating Noise to Sensitivity in Private Data Analysis

*Springer Berlin Heidelberg*



Duchi et al. (2013)

Local Privacy, Data Processing Inequalities, and Statistical Minimax Rates

*Advances in Neural Information Processing Systems*



Garcelon et al. (2021)

Local Differential Privacy for Regret Minimization in Reinforcement Learning

*Advances in Neural Information Processing Systems*



# References



Gajane et al. (2018)

Corrupt Bandits for Preserving Local Privacy

*Proceedings of Machine Learning Research*



Chen et al. (2018)

(Locally) Differentially Private Combinatorial Semi-Bandits

*International Conference on Machine Learning*



Abdadi et al. (2016)

Deep Learning with Differential Privacy

*Association for Computing Machinery*